

CS 160: Lab/Assignment 4

Due at 11:59PM on March 14, 2017

Numbers 1 and 2 use slide 26 in the openmp slides

1. Create a parallel version of the pi program using a parallel construct. Please use the OpenMP runtime library routine `omp_get_wtime()` to measure the execution time of the computational section in the program.

Please see `num1.c`

`num_steps = 100000`

Threads: 4

Execution time: 0.004686 seconds

2. Create a parallel version of the pi program using a loop construct. Your goal is to minimize the number changes made to the serial program. Please use the OpenMP runtime library routine `omp_get_wtime()` to measure the execution time of the computational section in the program.

Please see `num2.c`

`num_steps = 100000`

Threads: 4

Execution time: 0.010468 seconds

3. Parallelize the matrix multiplication program in the file `matmul.c` attached. Can you optimize the program by playing with how the loops are scheduled? Please use the OpenMP runtime library routine `omp_get_wtime()` to measure the execution time of the computational section in the program.

Yes, the program can be optimized for stronger hardware. If there is a CPU bottleneck, however, parallelization takes more time.

Execution time:

The sequential version of matrix multiplication costs 4.904458 seconds

The parallel version of matrix multiplication costs 1.901743 seconds

4. Does the following OpenMP code segment parallelize the for-loop correctly or not? Why?

```

int i, j, a[MAX];
j=1;
#pragma omp parallel for
for (i=0; i<MAX; i++) {
    j=j+2;
    a[i]=comp(j);
}

```

No, the following OpenMP code does not segment parallelize the for-loop correctly. Each thread should have different values for *j*, but in this case, they start out with the same value for *j* because it is defined as the integer 1 outside the OpenMP code segment. The whole point of parallelized loops is so they can safely execute in any order without loop-carried dependencies, but the above code does not do this.

5. Consider the following OpenMP program segment.

```

int a=1, b=2, c=3, d=4;
...
#pragma omp parallel for private(b), firstprivate(c)
lastprivate(d)
{
    ...
}

```

(a). Are *a*, *b*, *c*, and *d* local or shared in the parallel region?

The variable *a* is shared.

The variable *b* is local.

The variable *c* is local.

The variable *d* is local.

(b). What are their initial values inside the parallel region?

Initial Value of *a* is 1

Initial Value of *b* is garbage (not initialized)

Initial Value of *c* is 3

Initial Value of *d* is garbage (not initialized)

6. The goal of the following OpenMP program is to calculate π in parallel. Which variables are shared and which variables are private in the parallel region of the program? Identify and fix all bugs in the program.

```
#include <stdio.h>
#define MAX_THREADS 4
static long num_steps = 100000000;
double step;
int main ()
{
    int i, j;
    double pi, full_sum = 0.0;
    double start_time, run_time, x;
    step = 1.0/(double) num_steps;
    for(j=1; j<=MAX_THREADS; j++)
    {
        omp_set_num_threads(j);
        full_sum = 0.0;
        start_time = omp_get_wtime();
        #pragma omp parallel private(x), private(i)
        {
            int id = omp_get_thread_num();
            int numthreads = omp_get_num_threads();
            double partial_sum = 0;
            for (i=id; i< num_steps; i+=numthreads)
            {
                x = (i+0.5)*step;
                partial_sum = partial_sum + 4.0/(1.0+x*x);
            }
            full_sum += partial_sum;
        }
        pi = step * full_sum;
        run_time = omp_get_wtime() - start_time;
        printf("\n pi is %f in %f seconds %d threads \n ", pi,
run_time, j);
    }
}
```

In the code provided by the assignment...

Private variables: id, numthreads, partial_sum, i

Shared variables: x, full_sum